

Reproducibility Report: Baxter-King, Coppock, Straus and Vavreck (2025)

Endorsements vs. Information: Experimental Evidence of Backlash and Parallel Persuasion during the COVID-19 Public Health Crisis

Table of contents

1	Paper Overview	3
2	Summary	3
2.1	Does the deposited archive run?	4
2.2	Does the maintained rewrite reproduce the paper?	4
3	Original Archive Reproducibility	4
3.1	What runs, and what a stripped copy shows	4
3.2	The two scripts that fail	5
3.3	Checksums	5
4	Errata	5
5	Ground Truth Verification	6
5.1	Where the rewrite differs from the article	7
5.2	Coverage of the published floats	8
5.3	Every number in the article, extracted	8
6	Maintained Rewrite	9
7	Figure Verification	10
8	Maintained Rewrite Verification	13
9	R Environment	13

Drafted by Claude Opus 5 under the supervision of Alex Coppock.

This repository holds the actively maintained replication code for Baxter-King, Coppock, Straus and Vavreck (2025), together with the reproducibility report that documents what the original archive did and did not do. It is part of a program applying the maintenance proposal in Peer, Orr and Coppock (2021, *PS: Political Science & Politics*, doi [10.1017/S1049096521000366](https://doi.org/10.1017/S1049096521000366)) to a set of published archives.

Article	10.1093/pnasnexus/pgaf185
Replication archive	osf.io/bwv3c
Pre-analysis plans	None; the article states the analyses were not preregistered

The data are not redistributed here. The deposit lives on OSF and that is the only copy this repository points at. `download_original.R` fetches it and verifies every file; `original_manifest.csv` records the OSF file identifier, byte size, MD5 and SHA-256 of each of the 101 deposited files. OSF publishes two hashes of one set of stored bytes, computed by two different algorithms, so both are genuine checks and the script gates on both plus byte size. It also stops if `original/` holds anything the manifest does not list, and stops if a file that is already present no longer matches, which is how a mid-run write into the deposit would be caught rather than silently repaired. Re-hosting a copy would create a second archive that can drift from the first, which is the problem this project exists to document rather than to add to.

Repository layout. `maintained/` is the maintained rewrite: one script per published table or figure, writing to `output/`, which is committed so a reader can compare a fresh run against it without downloading anything. `ground_truth/` ties every published number to the code that produces it, and `maintained/in_text_claims.R` prints the same numbers again by an independent route. `original/` is created by the download script and is deliberately absent from the repository. This README is the reproducibility report, also available as a PDF in `report/`.

License. CC0 1.0 Universal, matching the terms of the deposit this repository maintains, so nothing in the chain is more restrictive than the archive itself. See `LICENSE`.

To reproduce. Clone or download the repository, open `baxter-king_etal_2025.Rproj`, and run:

```
source("run_all.R")
```

That fetches the deposited archive from OSF, verifies its checksums, fits every model, writes every table and figure into `maintained/output/`, runs the deposited code twice in a throwaway copy, rebuilds the ground truth, prints the in-text claims, and re-verifies the deposit. It takes about ten minutes. Individual scripts can be run on their own, but the three `analysis_*.R` scripts come before the tables and figures that read their model objects.

Required packages: tidyverse, estimatr, patchwork, modelsummary, knitr, kableExtra, scales, openssl, lmtest, here. Paths resolve through [here](#), so nothing depends on the working directory. A successful run overwrites `maintained/output/`, which is committed: **git diff on that folder is the reproduction check**, and the CSV, TeX, PNG and RDS output should come back byte-identical. The five PDF figures always show as changed, because a PDF records the time it was written; compare their PNG twins instead.

1 Paper Overview

Citation: Baxter-King, R., Coppock, A., Straus, G. and Vavreck, L. (2025). Endorsements vs. information: experimental evidence of backlash and parallel persuasion during the COVID-19 public health crisis. *PNAS Nexus*, 4(6), pgaf185.

Research question: Which of three kinds of persuasive appeal, endorsements, guidance and mandates, and factual information, move intentions to vaccinate and attitudes to masking, and do their effects differ by party?

Design: Ten survey experiments fielded in the UCLA COVID-19 Health and Politics Project between October 2020 and October 2022, among a total the article reports as 85,191 respondents. Two endorsement experiments randomize the identity of the person endorsing a vaccine; three guidance experiments randomize whether respondents are told the CDC recommends a masking policy or that vaccination is required for an activity; five information experiments randomize factual information about contagiousness and booster effectiveness. All estimates are weighted, use the Lin covariate adjustment, and take HC1 robust standard errors. Effects are reported overall and conditional on seven-point party identification.

Main finding: Endorsements by political figures polarize. Trump’s endorsement lowered intentions to vaccinate by 9.2 points on average in October 2020 while raising them among strong Republicans, and by April 2021 endorsements by Biden and Obama lowered intentions among the still unvaccinated in both parties. Guidance and information move Democrats and Republicans in parallel: telling respondents the CDC recommended a masking policy raised support for it by 6 points with no partisan interaction, and information about the bivalent booster raised booster intentions by 8.3 points, again without partisan heterogeneity.

2 Summary

The ground truth carries 484 rows. Where the deposited code produces a number the article prints, it reproduces it: 429 such comparisons agree at the precision the page prints and 1 do not. The maintained rewrite agrees with the article on 464 values and disagrees on 14, and 6 rows are unverifiable because the article states no number.

2.1 Does the deposited archive run?

The deposit ships 101 files of code, data and metadata and no output directory. Run in order, 61 of its 63 scripts complete and 2 fail. Every write in the deposit goes to a directory the deposit does not ship, so a full run overwrites 0 of the 101 deposited files. Running the archive in place would therefore be safe here, which is not something a reader could know without measuring it.

2.2 Does the maintained rewrite reproduce the paper?

Yes, with one class of exception. Every coefficient, standard error, sample size and R-squared in Tables 2 and 3 and in appendix Tables S1 to S4 reproduces, as does every estimate and standard error printed on Figures 1, 3 and 5 and every number the article states in prose or in a footnote. The 14 disagreements are the sample sizes of Table 1, which count respondents enrolled in each experiment while the deposited analysis files hold fewer rows, and one appendix claim that turns on a coding error in the deposit.

3 Original Archive Reproducibility

3.1 What runs, and what a stripped copy shows

The deposit is run twice in a throwaway copy, never in place. The first pass runs its scripts in the order `code/main.R` sources them and lets each see what the ones before it wrote. The second runs each script against a copy holding only deposited files, which is the test of whether a script passes only because another script left something behind.

Table 1: Deposited scripts, as shipped and against a copy holding only deposited files

Pass	Errors	Ran clean
as shipped	2	61
stripped	15	48

13 scripts run in the first pass and fail in the second. All of them read a model object an earlier script wrote, which is the deposit working as intended rather than a defect: `code/main.R` sources the analysis scripts before the tables and figures. It does mean that no table or figure script in this deposit can be checked on its own.

3.2 The two scripts that fail

Table 2: Deposited scripts that fail when the deposit is run in order

Script	Error
code/figures/figure_mini_meta.R	Error in eval(ei, envir) : object 'four_colors' not found
code/figures/appendix_figures_D24_D45.R	Error in 'method(update_ggplot, list(class_any, ggplot2::ggplo

`code/figures/figure_mini_meta.R`, which draws appendix Figure S62, never sources `code/helpers.R` and so cannot find the colour palette it uses. `code/figures/appendix_figures_D24_D45.R`, which draws appendix Figures S24 to S45, composes plots with + without loading `patchwork`; under `ggplot2 4.x` that is an error rather than the silent success it once was. Both are one-line fixes, and both mean that appendix figures the article cites cannot be redrawn from the deposit as it stands.

`code/figures/appendix_figures_E48_E59.R` is a third script worth naming: it exists in the deposit, the deposit's own README describes it, and `code/main.R` never sources it.

3.3 Checksums

All 101 deposited files match the MD5, the SHA-256 and the byte size that OSF publishes for them, and `original/` holds nothing else. The manifest records those published values, so a reader can verify a fresh download against the same numbers; the report does not quote a post-run hash of any file this pipeline writes, because a PDF records the time it was written and its hash changes on every run.

4 Errata

The maintained rewrite corrects one coding error in the deposit. It changes a claim the article makes.

The standard error of a difference in conditional average treatment effects. The deposit's `dic_estimator()` in `code/helpers.R` computes

```
std.error = sqrt(std.error 2 ^2 + std.error 2^2)
```

The second term should be `std.error_1`. As written, the Republican standard error is squared twice and the Democratic one never enters, so every standard error, p-value, confidence interval and equivalence-test p-value in the appendix’s equivalence tables is computed from the wrong quantity. This is a typographical error rather than an analytical choice, so the rewrite uses `sqrt(std.error_1^2 + std.error_2^2)`.

The correction moves a published claim. The article says that among the four vaccine mandate vignettes “we can affirm equivalence at 10 points in two cases but cannot in two cases.” Under the deposit’s standard error 2 of the four vignettes affirm equivalence at ten points; under the corrected one, 1 does. The ground truth records this row as a mismatch with `defect_locus = archive`.

Two further defects in the deposit are reported rather than corrected, because correcting them would change the published figures rather than a computation.

The series labels of Figure 2 are swapped. `code/analysis/run_models_w3_endorse_by_arm.R` fits the models named `*_prosocial` on the subset `experiment_arm == "A (Personal)"` and the models named `*_not_prosocial` on `experiment_arm == "B (Social)"`. `code/figures/figure_2.R` then relabels Prosocial as Social and Not Prosocial as Personal. Each framing’s estimates are therefore plotted under the other framing’s label. The rewrite labels each series by the `experiment_arm` value the model was fitted on.

The rows of the guidance equivalence table are swapped. `code/tables/equiv_test_guidance.R` maps the label G-1: CDC Mask Guidance 1 to the model object `w7_fit_cdcmask` and G-3: CDC Mask Guidance 2 to `w6_fit_cdcmask`. G-1 is the June 2021 experiment, whose analysis file is `w6`, and G-3 is the September 2021 experiment, whose file is `w7`. The two rows of the published table carry each other’s numbers.

5 Ground Truth Verification

`ground_truth/build_ground_truth.R` runs as the last step of `run_all.R`. Every published value in it was read from the article or its appendix and is used only as a comparison target; the archive column is read from a fresh run of the deposited code and the rewrite column from `maintained/output/`, so neither can drift from the code that made it. Published values are carried as the digits the page prints, and a value agrees when the pipeline’s number, printed to that precision, gives the same digits.

Table 3: Ground truth by location in the article

Table or figure	Rows	Archive agrees	Rewrite agrees	Rewrite differs
Appendix Table S2	80	80	80	0

Appendix Table S3	80	64	80	0
Appendix Table S1	70	70	70	0
Table 3	50	50	50	0
Appendix Table S4	40	32	40	0
Figure 1	32	32	32	0
Figure 3	32	32	32	0
Figure 5	32	32	32	0
Table 2	20	20	20	0
Results, endorsements	14	8	12	1
Table 1	10	0	0	10
Materials and methods	7	2	4	2
Results, information	5	5	5	0
Results, guidance	3	2	2	0
Figure 2	2	0	0	0
Footnote c	2	0	2	0
Abstract and Introduction	1	0	0	1
Appendix Figure S47	1	0	1	0
Figure 4	1	0	0	0
Footnote f	1	0	1	0
Footnote g	1	0	1	0

5.1 Where the rewrite differs from the article

Table 4: Every row where the maintained rewrite disagrees with the article

Claim	Article	Rewrite	Locus
E-1: Vaccine endorsement 1, N	14,946	14906	archive
E-2: Vaccine endorsement 2, N	7,249	7197	archive
G-1: CDC mask guidance 1, N	30,857	30751	archive
G-2: Vaccine mandate vignettes, N	10,298	10275	archive
G-3: CDC mask guidance 2, N	33,088	32943	archive
I-1: Contagiousness conversation, N	8,710	8661	archive
I-2: Delta variant conversation, N	8,710	8692	archive
I-3: Bivalent booster information, N	10,700	10677	archive
I-4: Bivalent booster information (children), N	1,628	1626	archive
I-5: Holiday surge information (children), N	1,715	1711	archive
Survey respondents across the ten experiments	85,191	96567	archive
E-2 unvaccinated respondents interviewed	7,249	7197	archive
Mandate vignettes affirming equivalence at 10 points	2	1	archive

10 of these are the N column of Table 1, which reports how many respondents were enrolled in each experiment. The deposited analysis file for each experiment holds fewer rows, and the deposit contains no script that assembles Table 1, so nothing in the archive can produce the published figures. One more is the headline count of respondents across the whole project: the union of the respondent identifiers in the ten deposited analysis files is larger than the 85,191 the abstract reports, so the deposit does not contain the file from which that number was counted. Another is the same gap restated in the text for E-2. The remaining 2 are the equivalence claim discussed above and one appendix claim about how often the interaction terms and the differences in conditional average treatment effects agree.

5.2 Coverage of the published floats

The float list is taken from the article and its appendix rather than from what the pipeline happens to produce: 74 floats in all, of which 13 carry at least one ground-truth row.

Table 5: Published floats, by whether the ground truth reaches them

Float group	Floats
Appendix section C, per-contrast panels	22
Appendix section D, predicted unvaccinated	23
Appendix section E, vignette CATEs	14
Meta-analysis of all studies	1
Verified cell by cell	13
Weighted against unweighted	1

The 61 uncovered floats are appendix figures. None prints a number: they are two-panel plots of group means and average causal effects, in the same form as Figures 1, 3 and 5, and the coefficients behind them are the ones appendix Tables S1 to S4 print, which the ground truth verifies cell by cell. Each uncovered float carries its reason in `ground_truth/float_coverage.csv`, and the build stops if any float has neither a row nor a reason.

5.3 Every number in the article, extracted

`ground_truth/published_claims.csv` is the exhaustive extraction: every numeric token in the article and its appendix, classified by hand.

Table 6: Numeric claims in the article and appendix, by type

Claim type	Claims
pipeline	476
structural	12
descriptive	8
transcribed	6
definitional	3

A pipeline or descriptive claim must have a row in the ground truth and a block in `maintained/in_text_claims.R`; the build stops if either is missing. `definitional`, `structural` and `transcribed` claims have no pipeline counterpart by construction: a scale endpoint is checked against the survey instrument, a count of experimental groups against the pipeline object with one row per group, and a number the article takes from another source against that source.

`maintained/in_text_claims.R` prints every one of these numbers beside the article's own sentence, reaching each by its own route from the same committed output the ground truth reads. Two independent derivations that disagree have found something.

6 Maintained Rewrite

The rewrite replaces 63 deposited scripts with 20 files. The deposit fits each model in its own script and repeats the same forty-line block for every endorser; the rewrite fits all of them by mapping one function over a named list of arms, so a change to the specification happens in one place.

Original	Replacement
<code>rm(list = ls())</code>	omitted
<code>setwd()</code>	<code>here::here()</code> with <code>.here</code> at the repository root
<code>library()</code> in each script	<code>source(here::here("maintained", "helpers.R"))</code>
<code>coefplot::position_dodgev()</code>	<code>position_dodge()</code> on the discrete y axis
<code>saveRDS()</code> repeated once per model	<code>map()</code> over the arms, then <code>walk2()</code>
Model objects saved with their environment attached	environment dropped, so two runs write identical bytes
<code>readRDS()</code> / <code>saveRDS()</code>	<code>read_rds()</code> / <code>write_rds()</code>


Original	Replacement
%>%	\ >
Tidy summaries saved where the table needs a fitted object	fitted objects saved throughout





Two substitutions are worth naming. `coefplot` has not been updated since 2018 and its `position_dodgev()` is built on `ggplot2` 3.x internals that emit a warning for every facet under `ggplot2` 4.x; `position_dodge()` handles a discrete y axis natively and the panels are equivalent. And the deposit saved the mandate vignette models as tidy tables of coefficients, which is why the deposit records no sample size or R-squared for appendix Tables S3 and S4; the rewrite saves the fitted objects, so those cells have a counterpart.

The rewrite also adds four table scripts the deposit lacks or leaves incomplete: appendix Tables S1 to S4, one script each, each writing its cells unrounded to a CSV alongside the formatted LaTeX. Nothing downstream has to read a value back out of a formatted table and round it twice.

7 Figure Verification


Figures 1, 3 and 5 print every estimate and standard error on the panel, in percentage points. All 96 of those numbers reproduce.







Figures 2 and 4 print no numbers. Each figure script writes a CSV of every estimate it plots, so the comparison is against the appendix regression tables that print the same coefficients rather than against the panel.

October 2020



April 2021





8 Maintained Rewrite Verification

Nothing in this pipeline draws at random. There is no seed anywhere in the rewrite because there is nothing to seed: every estimate is a weighted least squares fit and every test is analytic. Running `run_all.R` twice from a clean session returns byte-identical CSV, TeX, PNG and RDS output; only the five figure PDFs differ, because a PDF records the time it was written.

The acceptance test for a full run is that `git status --porcelain maintained/output/` shows only those five PDFs.

9 R Environment

Table 7: Package versions this report was built under

Component	Version
R	4.6.0
tidyverse	2.0.0
estimatr	1.0.6
modelssummary	2.6.0
ggplot2	4.0.3
patchwork	1.3.2
kableExtra	1.4.0
here	1.0.2